

SUBHOJYOTI MUKHERJEE

Wisconsin Institute of Discovery
University of Wisconsin-Madison
Madison, WI 53715

Email: smukherjee27@wisc.edu,
subhojyotimukherjee22@gmail.com
Website: <https://subhojyoti.github.io/>

Research Interests Reinforcement Learning ([RL](#)), Large Language Models ([LLM](#)), Reinforcement Learning with Human Feedback ([RLHF](#)), Incontext Learning ([ICL](#)), Optimal Design ([OD](#)).

Education **University of Wisconsin-Madison**, Madison, USA Fall 2019 – 2024 (expected)
Ph.D., Electrical & Computer Engineering
Adviser: Dr. Robert Nowak, Dr. Josiah Hanna and Dr. Qiaomin Xie

University of Wisconsin-Madison, Madison, USA Fall 2019 – 2021
M.S., Electrical Engineering
Adviser: Dr. Robert Nowak

Indian Institute of Technology Madras, India 2015–2018
M.S (Research), Computer Science
Advisers: Dr. Balaraman Ravindran and Dr. Nandan Sudarsanam

West Bengal University of Technology, Kolkata, India 2009–2013
Bachelor of Technology, Computer Science & Engineering

Publications

1. **Subhojyoti Mukherjee**, Anusha Lalitha, Kousha Kalantari, Aniket Anand Deshmukh, Ge Liu, Yifei Ma, Branislav Kveton, "*Optimal Design for Human Preference Elicitation*". (**NeurIPS 2024, main conference**) [Paper] ([LLM](#), [RLHF](#))
2. **Subhojyoti Mukherjee**, Josiah Hanna, Robert Nowak, "*SaVeR: Optimal Data Collection Strategy for Safe Policy Evaluation in Tabular MDP*". (**ICML 2024, main conference**) [Paper] ([RL](#))
3. **Subhojyoti Mukherjee**, Anusha Lalitha, Kousha Kalantari, Aniket Anand Deshmukh, Ge Liu, Yifei Ma, Branislav Kveton, "*Optimal Design for K-Way Human Feedback*". (**Models of Human Feedback for AI Alignment workshop ICML 2024**) [Paper] ([LLM](#), [RLHF](#), [OD](#))
4. Aniruddha Bhargava, Lalit Jain, Branislav Kveton, Ge Liu, **Subhojyoti Mukherjee**, "*Off-Policy Evaluation from Logged Human Feedback*". (**Models of Human Feedback for AI Alignment workshop ICML 2024**) [Paper] ([LLM](#))
5. **Subhojyoti Mukherjee**, Qiaomin Xie, Josiah Hanna, Robert Nowak, "*SPEED: Optimal Experimental Design for Policy Evaluation in Linear Heteroscedastic Bandits*". (**AISTATS 2024**) [Paper] ([RL](#), [OD](#))
6. **Subhojyoti Mukherjee**, Qiaomin Xie, Josiah Hanna, Robert Nowak, "*Multi-task Representation Learning for Pure Exploration in Bilinear Bandits*", Neural Information Processing Systems. (**NeurIPS 2023**) [Paper] ([RL](#), [OD](#))
7. **Subhojyoti Mukherjee**, Josiah Hanna, Robert Nowak, "*ReVar: Strengthening Policy Evaluation via Reduced Variance Sampling*". Uncertainty in Artificial Intelligence. (**UAI-22**) [Paper] ([RL](#))
8. **Subhojyoti Mukherjee**, "*Safety Aware Changepoint Detection for Piecewise i.i.d. Bandits*". Uncertainty in Artificial Intelligence (**UAI-22**). [Paper] ([RL](#))

9. **Subhojyoti Mukherjee***, Ardhendu Tripathy*, Robert Nowak, "*Chernoff Sampling for Active Testing and Extension to Active Regression*". The 25th International Conference on Artificial Intelligence and Statistics (**AISTATS-22**). [Paper] ([RL](#), [OD](#))
10. Blake Mason, Romain Camilleri, **Subhojyoti Mukherjee**, Kevin Jamieson, Robert Nowak, Lalit Jain, "*Nearly Optimal Algorithms for Level Set Estimation*". The 25th International Conference on Artificial Intelligence and Statistics (**AISTATS-22**). [Paper] ([RL](#), [OD](#))
11. Samarth Gupta, Shreyas Chaudhari, **Subhojyoti Mukherjee**, Gauri Joshi, Osman Yagan, "*A Unified Approach to Translate Classical Bandit Algorithms to the Structured Bandit Setting*", *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP-21)*. [Paper] ([RL](#))
12. Samarth Gupta, Shreyas Chaudhari, **Subhojyoti Mukherjee**, Gauri Joshi, Osman Yagan, "*A Unified Approach to Translate Classical Bandit Algorithms to the Structured Bandit Setting*", *IEEE Journal on Selected Areas in Information Theory* (**2020**). [Paper] ([RL](#))
13. **Subhojyoti Mukherjee**, and Odalric-Ambrym-Maillard, "*Distribution-dependent and Time-uniform Bounds for Piecewise i.i.d Bandits*", *Thirty-sixth International Conference on Machine Learning (ICML-19)*, Workshop on Reinforcement Learning for Real Life 2019 track [Poster]. [Paper] ([RL](#))
14. **Subhojyoti Mukherjee**, K.P. Naveen, Nandan Sudarsanam, and Balaraman Ravindran, "*Efficient UCBV: An Almost Optimal Algorithm using Variance Estimates*", *Proceedings of the Thirty-Second Association for the Advancement of Artificial Intelligence (AAAI-18)*, main conference track [Oral]. [Paper] ([RL](#))
15. **Subhojyoti Mukherjee**, K.P. Naveen, Nandan Sudarsanam, and Balaraman Ravindran, "*Thresholding Bandits with Augmented UCB*", *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI-17)*, main conference track [Poster]. [Paper] ([RL](#))

Preprints

1. **Subhojyoti Mukherjee**, Anusha Lalitha, Sailik Sengupta, Aniket Anand Deshmukh, Branislav Kveton, "*Multi-Objective Alignment of Large Language Models Through Hypervolume Maximization*". AISTATS 2025 (Submitted) [Paper] ([LLM](#), [RLHF](#))
2. **Subhojyoti Mukherjee**, Ge Liu, Aniket Anand Deshmukh, Anusha Lalitha, Yifei Ma, Branislav Kveton, "*Optimal Design for Adaptive In-Context Prompt Tuning in Large Language Models*". NeurIPS 2024 (Submitted) [Paper] ([LLM](#), [ICL](#), [OD](#))
3. **Subhojyoti Mukherjee**, Josiah Hanna, Qiaomin Xie, Robert Nowak, "*Pretraining Decision Transformers with Reward Prediction for In-Context Structured Bandit Learning*". NeurIPS 2024 (Submitted) [Paper] ([LLM](#), [ICL](#))
4. **Subhojyoti Mukherjee**, Ruihao Zhu, Branislav Kveton, "*Efficient and Interpretable Bandit Algorithms*", [Paper]. ([RL](#), [OD](#))
5. **Subhojyoti Mukherjee**, Devin Conathan, Robert Nowak, "*AdaTune: Active Learning for Fine-Tuning BERT on QA Task*" ([LLM](#), [RL](#))

Research Internships	<ol style="list-style-type: none"> Amazon AWS AI, Santa Clara, USA: Summer 2024 (Full-time), Host: Branislav Kveton, Yifei Ma, Anusha Lalitha, Kousha Kalantiri, Aniket Deshmukh. Working on Alignment for Multi-objective optimization with LLMs. Amazon AWS AI, Santa Clara, USA: Fall 2023 (Part-time), Host: Branislav Kveton, Yifei Ma, Anusha Lalitha, Kousha Kalantiri, Ge Liu, Aniket Deshmukh, Anoop Deoras. Working on RLHF with LLMs Amazon AWS AI, Santa Clara, USA: Summer 2023 (Full-time), Host: Branislav Kveton, Yifei Ma, Anusha Lalitha, Ge Liu, Aniket Deshmukh, Anoop Deoras. Worked on Active In-Context Learning with LLMs CMU, ECE Dept., Pittsburgh, USA: Summer 2019, Host: Gauri Joshi. Worked on Structured Bandits Adobe Research, San Jose, USA: Spring 2018. Host: Branislav Kveton. Worked on Item recommendation with Ranking and Bandits INRIA, SequeL Lab, Lille, France: Fall 2017, Host: Odalric Maillard. Worked on Non-stationary Bandits 	
Master's Thesis (EE, UW-Madison)	Active Sequential Hypothesis Testing with Extension to Active Regression and Multi-armed Bandits [Thesis] (RL , OD)	
Master's Thesis (CS, IIT Madras)	Finite-time Analysis of Frequentist Strategies for Multi-armed Bandits [Thesis](RL)	
Teaching Experience	Teaching Assistant , UW-Madison <i>Matrix Methods in Machine Learning</i> - Prof. Robert Nowak <i>Mathematical Foundation in Machine Learning</i> - Prof. Robert Nowak Teaching Assistant , UMass Amherst <i>Natural Language Processing</i> - Prof. Mohit Iyyer <i>Design of Algorithms</i> - Prof. Daniel Sheldon Teaching Assistant , IIT Madras <i>Introduction to Programming</i> - Prof. Raghavendra Rao B. V. <i>Reinforcement Learning</i> (twice) - Prof. Balaraman Ravindran	2019–current 2018–2019 2015–2018
Reviewer and Service	<ol style="list-style-type: none"> AISTATS, UAI, AAAI, ICML, ICLR, NeurIPS, TMLR, KDD, RLC Main Co-ordinator of SILO seminar at UW-Madison 	
Award Grants and Fellowship	<ol style="list-style-type: none"> Top reviewer award for UAI 2023, Neurips 2023 Student Scholarship for AAAI 2018, UAI 2022, Neurips 2023 UW-Madison nominee for Apple PhD fellowship and Two-sigma PhD fellowship. Received UW-Madison Chancellor's Opportunity Fellowship 2019-20, UW-Madison ECE Welcome Award of USD 3000. IIT Madras student travel grant of USD 2300, Google travel grant of USD 1700, Microsoft travel grant of USD 1435 (declined). 	
Other Achievements	Ranked 1150/155190 candidates in Graduate Aptitude Test in Engineering (GATE) 2014. Secured 98.93 percentile in Common Admission Test (CAT) 2014 among 196988 candidates.	
References	Available Upon Request.	